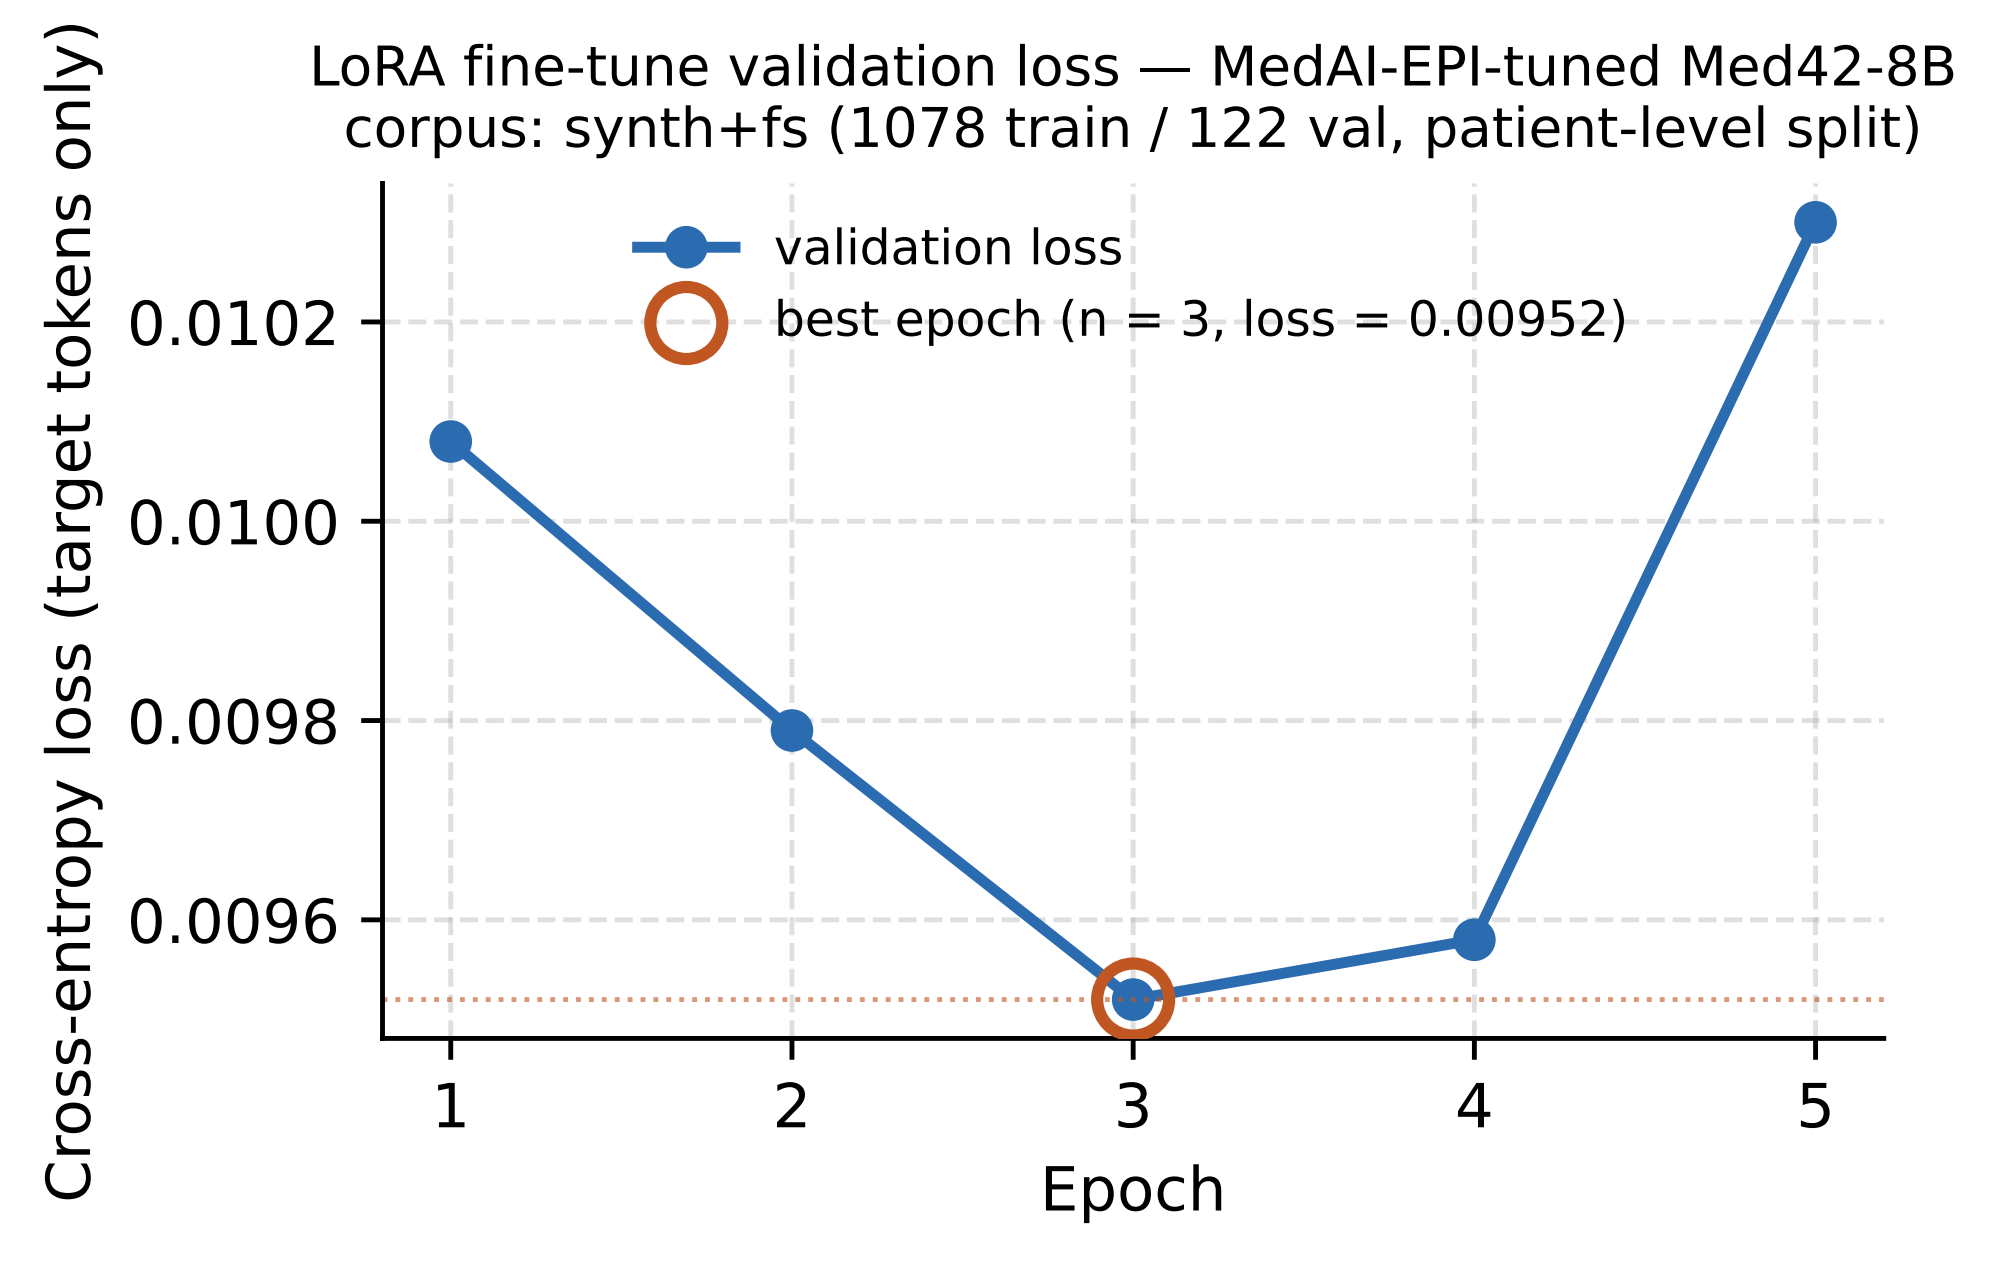
LoRA fine-tune validation loss — MedAI-EPI-tuned Med42-8B
corpus: synth+fs (1078 train / 122 val, patient-level split)
validation loss
best epoch (n = 3, loss = 0.00952)
Cross-entropy loss (target tokens only)
0.0102
0.0100
0.0098
0.0096
1
2
3
4
5
Epoch